

Design and Implementation of a Single-Cycle MIPS32 Processor

**A PROJECT REPORT OF
COMP-261: Computer Organization and Assembly Language**

SUBMITTED BY:

Areeba Shafqat (B24F0209CS187)

Hafsa Aftab (B24F0247CS217)

Urooj Azmat (B24F2067CS201)

*in partial fulfillment for the award of the degree
of*

Bachelor of Computer Science 2024-2028



SUBMITTED TO:

Instructor: Dr. Muhammad Ejaz Khan

Lab Engineer: Ms. Kouser Bibi

**Pak-Austria Fachhochschule:
Institute of Applied Sciences and Technology Mang, Haripur
(Spring 2026)**

ABSTRACT

This report presents the design, implementation, and verification of a simplified single-cycle 32-bit MIPS processor for COMP-261 (Computer Organization and Assembly Language) at Pak-Austria Fachhochschule (PAF-IAST). The processor incorporates a hardwired control unit and a single-cycle datapath comprising a Program Counter (PC), Instruction Memory, Register File (32×32-bit with two read ports and one write port), Arithmetic Logic Unit (ALU) supporting add, sub, slt, and, or operations, Data Memory, Sign-Extend unit, and necessary multiplexers. The control unit generates eight control signals (RegDst, ALUSrc, MemtoReg, RegWrite, MemRead, MemWrite, Branch, ALUOp) based on the opcode field, with a two-level ALU decoder for R-type instructions using the funct field.

A twelve-instruction MIPS32 assembly test program was written to demonstrate all required instruction types: R-type (add, sub, slt), I-type memory instructions (lw, sw), conditional branch (beq), and immediate arithmetic (addi). The program operates on two pre-initialized operands (15 and 7), computes sum (22) and difference (8), performs a set-less-than comparison (1), stores the result to memory, reloads it for verification, and executes a taken branch that bypasses a NOP instruction. All instructions were converted to 32-bit hexadecimal machine code, manually verified using MIPS32 encoding rules, and simulated using the MARS simulator.

The report includes instruction format diagrams, a complete control signal truth table, Datapath block diagram, assembly program listing with hex code table, step-by-step execution trace, branch behavior analysis, simulation screenshots, and six trade-off analyses comparing single-cycle vs pipelined architectures, ALU complexity vs power dissipation, Harvard vs unified memory, hardwired vs microprogrammed control, early vs late branch resolution, and register file read ports. All technical content is grounded in Harris and Harris, Digital Design and Computer Architecture, Chapter 7.

Keywords: MIPS32, single-cycle processor, hardwired control unit, Datapath, RISC architecture, MARS simulator, computer organization.

TABLE OF CONTENTS

1 INTRODUCTION	5
2 MIPS INSTRUCTION FORMAT	5
2.1 R-TYPE FORMAT	6
2.2 I-TYPE FORMAT	7
2.3 INSTRUCTION ENCODING REFERENCE.....	7
3 DATAPATH DESIGN	8
3.1 PROGRAM COUNTER (PC).....	9
3.2 INSTRUCTION MEMORY	9
3.3 REGISTER FILE	9
3.4 SIGN-EXTEND UNIT.....	10
3.5 ALU AND ALU CONTROL	10
3.6 DATA MEMORY	10
3.7 MULTIPLEXERS	10
3.8 DATAPATH TRACE FOR KEY INSTRUCTIONS	11
4 CONTROL UNIT DESIGN.....	13
4.1 MAIN CONTROL UNIT TRUTH TABLE	13
4.2 ALU CONTROL SUB-UNIT	15
5 ASSEMBLY PROGRAM AND MACHINE CODE	15
5.1 COMPLETE COMMENTED ASSEMBLY LISTING	16
5.2 MACHINE CODE HEX TABLE.....	16
5.3 MANUAL BINARY-TO-HEX DERIVATION	17
5.3.1 R-type: add \$t2, \$t0, \$t1 (Instruction 3)	17
5.3.2 I-type: lw \$t0, 0(\$s0) (Instruction 1)	18
5.3.3 I-type (Branch): beq \$t4, \$t4, skip (Instruction 9)	18
6 SIMULATION AND VERIFICATION	18

6.1 INSTRUCTION # 01: lw \$s0, var_a.....	18
6.2 INSTRUCTION # 02: lw \$s1, var_b	19
6.3 INSTRUCTION # 03 (add s2 ,s0, \$s1)	19
6.4 INSTRUCTION # 04 (sub s3 ,s0, \$s1).....	19
6.5 INSTRUCTION # 05 (slt t4 ,s3, \$s2).....	20
6.6 INSTRUCTION # 06 (beq t4 ,t4, skip)	20
6.7 INSTRUCTION # 07 (addi zero ,zero, 0) SKIPPED	20
6.8 INSTRUCTION # 08 (sw \$s2, final)	21
6.9 INSTRUCTION # 09 (li \$v0, 10)	21
6.10 INSTRUCTION # 10 (syscall).....	22
6.11 SIMULATION SCREENSHOTS:	22
6.12 MEMORY CONTENTS	23
6.13 STEP-BY-STEP EXECUTION TRACE.....	23
6.14 BRANCH BEHAVIOUR ANALYSIS	24
6.15 FINAL STATE SUMMARY	24
6.16 MARS SIMULATION PROCEDURE	24
7 TRADE-OFF ANALYSIS: PERFORMANCE VS POWER.....	25
7.1 ALU COMPLEXITY VS DYNAMIC POWER ($P = ACV^2F$)	26
7.2 SEPARATE VS UNIFIED MEMORY	26
7.3 HARDWIRED VS MICROPROGRAMMED CONTROL	27
7.4 LATE VS EARLY BRANCH RESOLUTION.....	27
7.5 REGISTER FILE READ PORTS	27
8 CONCLUSION.....	27
9 REFERENCES	28

1 INTRODUCTION

The single-cycle MIPS32 processor represents the foundational microarchitectural paradigm in which every instruction, regardless of complexity, completes its entire execution within a single clock period. As defined by Harris and Harris in Chapter 7 of *Digital Design and Computer Architecture*, the single-cycle implementation fetches an instruction, decodes it, reads operands from the register file, executes an arithmetic or logical operation in the ALU, optionally accesses data memory, and writes a result back to the register file all within one uninterrupted clock cycle. The critical path length is determined by the longest instruction type, which in MIPS32 is the load word (lw) instruction, since it traverses the maximum number of Datapath stages.

This project fulfils the Problem Based Learning (PBL) requirements for COMP-261 as specified by Dr. Muhammad Ejaz Khan. The objective is to design, implement, and verify a simplified 32-bit MIPS processor incorporating a hardwired control unit and a single-cycle Datapath. The implementation is performed entirely through simulation using the MARS (MIPS Assembler and Runtime Simulator) environment, which faithfully models the architectural behavior of the physical hardware at the register-transfer level.

The instruction subset selected for this project encompasses all mandatory instruction types identified in the PBL specification: R-type instructions (add, sub, slt) which exercise the ALU directly with two register operands; I-type memory instructions (lw, sw) which demonstrate address computation and memory access paths; the conditional branch instruction (beq) which exercises the PC update logic and the Zero flag of the ALU; and the immediate arithmetic instruction (addi) which illustrates the ALUSrc multiplexer selecting the sign-extended immediate value. Together, these instructions constitute a representative and complete test of all major datapath components and control signal combinations.

The program operates on two pre-initialized operands stored in the .data segment: operandA = 15 (0x0000000F) and operandB = 7 (0x00000007). The processor performs addition, subtraction, a set-less-than comparison, stores the result to memory, reloads it for verification, and conditionally branches to a terminal routine. This sequence exercises every significant path through the single cycle Datapath as described in Harris and Harris, Chapter 7, Figures 7.13 through 7.27.

2 MIPS INSTRUCTION FORMAT

The MIPS32 instruction set architecture specifies a fixed 32-bit instruction width. All instructions must conform to one of three canonical formats: R-type (register), I-type (immediate), and J-type (jump). This project employs R-type and I-type formats exclusively, as described in Harris and Harris, Chapter 7, Section 7.1.

2.1 R-TYPE FORMAT


R-type instructions encode operations that use two source registers and write a result to a destination register. The 32-bit word is partitioned into six fields as follows:

op [31:26]	rs [25:21]	rt [20:16]	rd [15:11]	shamt [10:6]	funct [5:0]
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

Field definitions:

- op[31:26]: 6-bit opcode; always 000000 for all R-type instructions
- rs[25:21]: 5-bit first source register address
- rt[20:16]: 5-bit second source register address
- rd[15:11]: 5-bit destination register address
- shamt[10:6]: 5-bit shift amount; zero for non-shift operations
- funct[5:0]: 6-bit function code that selects the specific ALU operation

Datapath: Reg-Reg ALU Instructions



2.2 I-TYPE FORMAT

I-type instructions encode operations with one register source and a 16-bit signed immediate value. The four fields are:

op [31:26]	rs [25:21]	rt [20:16]	immediate [15:0]
6 bits	5 bits	5 bits	16 bits

Field definitions:

- op [31:26]: 6-bit opcode uniquely identifies instruction type
- rs [25:21]: 5-bit base register for memory address computation
- rt [20:16]: 5-bit destination register (lw/addi) or source (sw/beq)
- immediate [15:0]: 16-bit value sign-extended to 32 bits

Datapath: Reg-Imm ALU Instructions




Figure 2.2: Datapath for I-type (Register-Immediate) ALU Instructions

2.3 INSTRUCTION ENCODING REFERENCE

The following table documents the precise binary encoding of every instruction employed in the test program, providing opcode values, register numbers in binary, and function codes:

Instruction	Type	op [31:26]	rs [25:21]	rt [20:16]	rd [15:11]	shamt	funct / imm
add	R	000000	01000 (8)	01001 (9)	01010 (10)	00000	100000

sub	R	000000	01000 (8)	01001 (9)	01011 (11)	00000	100010
slt	R	000000	01011 (11)	01010 (10)	—	00000	101010
lw	I	100011	10000 (16)	01000 (8)	N/A	N/A	offset = 0
sw	I	101011	10000 (16)	01010 (10)	N/A	N/A	offset = 8
beq	I	000100	01100 (12)	01100 (12)	N/A	N/A	offset = 1
addi	I	001000	01010 (10)	01010 (10)	N/A	N/A	imm = 0

MIPS Instruction Formats




Figure 2.1: MIPS Instruction Format Field Definitions

3 DATAPATH DESIGN

The single-cycle MIPS Datapath, as illustrated in Harris and Harris Chapter 7 Figure 7.14, interconnects the following major components through a network of datapaths and multiplexers, all orchestrated by the hardwired control unit. Each component is described below in terms of its function, interface, and role in instruction execution. The single-cycle MIPS datapath interconnects all major components as shown in the block diagrams below.




Figure 3.1: Complete single-cycle MIPS processor Datapath with control unit

3.1 PROGRAM COUNTER (PC)

The Program Counter is a 32-bit register that holds the address of the currently executing instruction. At the rising edge of each clock cycle, the PC is updated either to PC+4 (sequential execution) or to the branch target address $PC+4+(\text{sign_extend}(\text{offset})\ll 2)$, depending on the PCSrc control signal. PCSrc is asserted only when both the Branch control signal and the ALU Zero flag are simultaneously active. In MARS simulation, the PC is initialized to the first instruction in the .text segment at address 0x00400000.

3.2 INSTRUCTION MEMORY

The Instruction Memory is a read-only combinational memory, indexed by the PC value, that outputs a 32-bit instruction word. In the single-cycle model it is accessed purely combinational there is no clock edge required for the read (Harris Ch.7, Section 7.3.1). In this project the instruction memory is pre-loaded with the twelve instructions of the test program. The separate Harvard architecture approach (distinct Instruction and Data Memory modules) eliminates structural hazards that would arise if a unified memory were shared for instruction fetch and data access simultaneously.

3.3 REGISTER FILE

The MIPS Register File contains 32 general-purpose 32-bit registers, addressed by 5-bit register identifiers. It provides two simultaneous combinational read ports (RD1 and RD2, driven by rs and rt fields respectively) and one synchronous write port (WD3, driven at the clock edge when RegWrite is asserted). As specified in Harris Ch.7, the write is clocked while reads are asynchronous, enabling operand values to be available to the ALU within the same clock period. Register \$zero (register 0) is hardwired to zero and cannot be written.

3.4 SIGN-EXTEND UNIT

The Sign-Extend unit takes the 16-bit immediate field from the instruction [15:0] and produces a 32-bit value by replicating the most significant bit (bit 15) across the upper 16 positions. This preserves the two's complement signed value of the immediate field. The sign-extended value is then presented to the ALUSrc multiplexer. For lw and sw, it represents the memory offset; for beq, it represents the branch offset (in words, subsequently shifted left by 2); for addi, it represents the addend.

3.5 ALU AND ALU CONTROL

The Arithmetic-Logic Unit performs the core computation for all instruction types. It receives two 32-bit operands: SrcA (always from the register file RD1 port) and SrcB (selected by the ALUSrc multiplexer from either RD2 or the sign-extended immediate). The ALU supports addition, subtraction, bitwise AND, bitwise OR, and set-less-than (SLT) operations. It produces a 32-bit result and a 1-bit Zero flag, which is asserted when the result is exactly zero. The Zero flag is the mechanism by which beq determines whether the two operands are equal. The ALU Control sub-unit receives ALUOp[1:0] from the main control unit and, when ALUOp=10, also inspects funct[5:0] from the instruction word to determine the precise operation (Harris Ch.7, Figure 7.18, Table 7.3).

3.6 DATA MEMORY

The Data Memory is a word-addressable read/write memory. It is written when MemWrite is asserted (sw instruction) and read when MemRead is asserted (lw instruction). The address is the 32-bit result from the ALU (base register + sign-extended offset). The write data is sourced from register file read port RD2. The read data is presented to the MemtoReg multiplexer, which selects between the ALU result and the memory read data for the register file write-back.

3.7 MULTIPLEXERS

Three multiplexers implement the datapath routing decisions made by the control unit: (1) RegDst MUX selects either rt[20:16] (I-type) or rd[15:11] (R-type) as the destination register address for the write port of the register file. (2) ALUSrc MUX selects either RD2 from the register file or the sign-extended immediate as the second ALU operand. (3) MemtoReg MUX selects either the ALU result or the Data Memory read data as the value to write back to the register file. Additionally, the PCSrc MUX selects between PC+4 and the branch target address for the next PC value.

3.8 DATAPATH TRACE FOR KEY INSTRUCTIONS

R-type (add \$t2, \$t0, \$t1):

PC provides address → Instruction Memory returns 32-bit word → Control unit decodes $op=000000$, asserts $RegDst=1$, $ALUSrc=0$, $MemtoReg=0$, $RegWrite=1$, $ALUOp=10$ → Register File reads \$t0 (rs) and \$t1 (rt) → $ALUSrc$ MUX routes RD2 to ALU → ALU Control decodes $funct=100000$, performs addition → ALU result (25) → $MemtoReg$ MUX routes result → $RegDst$ MUX selects rd (\$t2) → Register File writes 25 to \$t2.




Figure 3.8.1: Datapath diagram for ADD Instruction showing active path

lw \$t0, 0(\$s0):

Control asserts $MemRead=1$, $ALUSrc=1$, $MemtoReg=1$, $RegWrite=1$ → RF reads \$s0 → $Sign-Extend(0)$ → $ALUSrc$ MUX routes imm to ALU → ALU adds base+offset → Data Memory read at computed address → value (15) routed through $MemtoReg$ MUX → written to \$t0.




Figure 3.8.2: Active Datapath for LW instruction

sw \$t2, 8(\$s0):

Control asserts $MemWrite=1$, $ALUSrc=1$, $RegWrite=0$ → RF reads \$s0 (for address) and \$t2 (write data) → ALU computes address → Data Memory written with \$t2's value.




Figure 3.8.3: Active Datapath for SW instruction

beq \$t4, \$t4, skip:

Control asserts Branch=1, ALUSrc=0, ALUOp=01 → RF reads \$t4 twice → ALU subtracts: $\$t4 - \$t4 = 0$ → Zero=1 → Branch AND Zero → PCsrc=1 → $PC \leftarrow PC + 4 + (1 \ll 2)$ = branch target (skip label). The beq is TAKEN.




Figure 3.8.4: Active Datapath for BEQ instruction



Figure 3.2: Summary of MIPS single-cycle Datapath data flow sequence

4 CONTROL UNIT DESIGN

As described in Harris and Harris Chapter 7 Section 7.3, the hardwired control unit implements purely combinational logic that maps the 6-bit opcode field `op [31:26]` of the fetched instruction to a set of control signals. For R-type instructions (opcode 000000), a secondary ALU Control sub-unit further decodes the `funct` field to determine the specific ALU operation. This two-level decoding strategy minimizes the number of signals that the main control unit must generate while supporting the full complement of R-type operations.

4.1 MAIN CONTROL UNIT TRUTH TABLE


The following table presents the complete control signal specification for every instruction class implemented in this project. An 'X' denotes a don't care condition where the signal value does not affect the outcome for that instruction:

Instruction	RegDst	ALUSrc	MemtoReg	RegWrite	MemRead	MemWrite	Branch	ALUOp [1:0]
R-type (add, sub, slt)	1	0	0	1	0	0	0	10
lw	0	1	1	1	1	0	0	00

sw	X	1	X	0	0	1	0	00
beq	X	0	X	0	0	0	1	01
addi	0	1	0	1	0	0	0	00

Signal definitions:

RegDst: selects rd (R-type) or rt (I-type) as write register. ALUSrc: selects immediate (1) or register RD2 (0) as second ALU operand. MemtoReg: routes memory data (1) or ALU result (0) to register file write port. RegWrite: enables write to register file. MemRead: enables data memory read. MemWrite: enables data memory write. Branch: enables conditional PC branch. ALUOp[1:0]: 00=add for address, 01=subtract for beq, 10=R-type determined by funct.






Figure 4:1 Single-cycle MIPS Datapath with hardwired control unit

4.2 ALU CONTROL SUB-UNIT

When the main control unit generates $ALUOp=10$ (R-type), the ALU Control sub-unit inspects the funct field to determine the operation. The complete decode table is as follows:


ALUOp [1:0]	funct [5:0]	ALU Operation	Used By
00	XXXXXX	ADD	lw, sw, addi
01	XXXXXX	SUB (for Zero flag)	beq
10	100000	ADD	add
10	100010	SUB	sub
10	101010	SLT	slt

This two-level control hierarchy, as illustrated in Harris Ch.7 Figure 7.18, minimises the complexity of the main control decoder. The ALU control sub-unit adds minimal gate delay and is far simpler than incorporating function-field decoding directly into the main control logic. The ALU Control sub-unit generates the 3-bit ALU Control signal (000=AND, 001=OR, 010=ADD, 110=SUB, 111=SLT), which directly drives the ALU operation selection mux.

5 ASSEMBLY PROGRAM AND MACHINE CODE

5.1 COMPLETE COMMENTED ASSEMBLY LISTING

The test program is structured in two segments consistent with MIPS convention. The .data segment preinitializes three words in Data Memory. The .text segment contains twelve instructions that exercise all required datapath paths:



```

1  .data
2  var_a:  .word 15
3  var_b:  .word 7
4  final:  .word 0
5  .text
6  main:
7  # 1. Load from memory
8  lw $s0, var_a      # $s0 = 15
9  lw $s1, var_b      # $s1 = 7
10 # 2. Arithmetic Operations
11 add $s2, $s0, $s1   # $s2 = 15 + 7 = 22
12 sub $s3, $s0, $s1   # $s3 = 15 - 7 = 8
13 # 3. Set Less Than (compare)
14 slt $t4, $s3, $s2   # $t4 = 1 (because 8 < 22)
15 # 4. Conditional Branch (taken)
16 beq $t4, $t4, skip  # branch taken to skip
17 # 5. NOP (skipped by branch)
18 addi $zero, $zero, 0 # skipped
19 skip:
20 # 6. Store result back to memory
21 sw $s2, final       # store 22 in memory
22 # 7. Terminate
23 li $v0, 10          # exit code
24 syscall

```

Figure 5.1: Complete MIPS32 assembly test program

Program Memory Map (.data segment, base ≈ 0x10010000):

- offset +0: operandA = 15 (0x0000000F)
- offset +4: operandB = 7 (0x00000007)
- offset +8: result = 0 (written by sw instruction, becomes 22)

5.2 MACHINE CODE HEX TABLE

The following table provides the 32-bit hexadecimal encoding for every instruction in the test program, derived by applying the MIPS32 encoding rules to each assembly instruction:

Sr	Assembly Instruction	32-bit Hex Code	Type	Opcode / Funct	Description
1	lw \$t0, 0(\$s0)	0x8E080000	I	op = 0x23	Load operand A (15) into \$t0
2	lw \$t1, 4(\$s0)	0x8E090004	I	op = 0x23	Load operand B (7) into \$t1
3	add \$t2, \$t0, \$t1	0x01095020	R	op = 0, fn = 0x20	Add \$t0 + \$t1 = 22 → \$t2
4	sub \$t3, \$t0, \$t1	0x01095822	R	op = 0, fn = 0x22	Subtract → \$t3 = 8
5	addi \$t2, \$t2, 0	0x214A0000	I	op = 0x08	Add immediate (format demo)
6	sw \$t2, 8(\$s0)	0xAE0A0008	I	op = 0x2B	Store 22 into memory
7	lw \$t5, 8(\$s0)	0x8E0D0008	I	op = 0x23	Reload result
8	slt \$t4, \$t3, \$t2	0x016A602A	R	op = 0, fn = 0x2A	Set \$t4 = 1 (8 < 22)
9	beq \$t4, \$t4, skip	0x118C0001	I	op = 0x04	Branch taken (equal)
10	addi \$zero, \$zero, 0	0x20000000	I	op = 0x08	NOP (skipped)
11	addi \$v0, \$zero, 10	0x2002000A	I	op = 0x08	Load exit code
12	syscall	0x0000000C	R	op = 0, fn = 0x0C	Program terminate

5.3 MANUAL BINARY-TO-HEX DERIVATION

The following subsections present the complete step-by-step binary encoding for three representative instructions, demonstrating the mechanical application of MIPS instruction formats.

5.3.1 R-type: add \$t2, \$t0, \$t1 (Instruction 3)

Register numbers: \$t0=8 (01000), \$t1=9 (01001), \$t2=10 (01010).

- op[31:26] = 000000 (R-type opcode)
- rs[25:21] = 01000 (\$t0 = register 8)
- rt[20:16] = 01001 (\$t1 = register 9)

- $rd[15:11] = 01010$ (\$t2 = register 10)
- $shamt[10:6] = 00000$ (no shift)
- $funct[5:0] = 100000$ (ADD function code)

Concatenated binary: 0000 0001 0000 1001 0101 0000 0010 0000

Grouped as hex nibbles: 0x01095020

5.3.2 I-type: lw \$t0, 0(\$s0) (Instruction 1)

Register numbers: \$s0=16 (10000), \$t0=8 (01000). Offset = 0 (0000000000000000).

- $op[31:26] = 100011$ (lw opcode = 35)
- $rs[25:21] = 10000$ (\$s0 = register 16)
- $rt[20:16] = 01000$ (\$t0 = register 8)
- $immediate[15:0] = 0000000000000000$ (offset = 0)

Concatenated binary: 1000 1110 0000 1000 0000 0000 0000 0000

Grouped as hex nibbles: 0x8E080000

5.3.3 I-type (Branch): beq \$t4, \$t4, skip (Instruction 9)

Register numbers: \$t4=12 (01100). Branch offset = +1 word (the beq skips the next addi instruction, so offset = 1 in words).

- $op[31:26] = 000100$ (beq opcode = 4)
- $rs[25:21] = 01100$ (\$t4 = register 12)
- $rt[20:16] = 01100$ (\$t4 = register 12)
- $immediate[15:0] = 0000000000000001$ (offset = +1 word ahead)

Concatenated binary: 0001 0001 0000 1100 0000 0000 0000 0001

Grouped as hex nibbles: 0x118C0001

Branch target address = $(PC+4) + (1 \times 4)$ = address of 'skip' label.

6 SIMULATION AND VERIFICATION

6.1 INSTRUCTION # 01: lw \$s0, var_a

Text segment after Execution:

Address	Code	Basic	
0x00400000	0x3c011001	lui \$1,0x00001001	8: lw \$s0, var_a # \$s0 = 15
0x00400004	0x8c300000	lw \$16,0x00000000(\$1)	

Registers:

\$s0	15	0x0000000f
\$16	16	0x0000000f

Data Segment:

Data Segment	
Address	Value (+0)
0x10010000	0x0000000f

6.2 INSTRUCTION # 02: lw \$s1, var_b

Text segment after Execution:

0x00400004	0x8c300000	lw \$16,0x00000000(\$1)	
0x00400008	0x3c011001	lui \$1,0x00001001	9: lw \$s1, var_b # \$s1 = 7

Registers:

\$s0	16	0x0000000f
\$s1	17	0x00000007

Data Segment:

Data Segment		
Address	Value (+0)	Value (+4)
0x10010000	0x0000000f	0x00000007

6.3 INSTRUCTION # 03 (add \$s2,\$s0,\$s1)

Text segment after Execution:

0x00400010	0x02119020	add \$18,\$16,\$17	11: add \$s2, \$s0, \$s1 # \$s2 = 15 + 7 = 22
0x00400014	0x02119822	sub \$19,\$16,\$17	12: sub \$s3, \$s0, \$s1 # \$s3 = 15 - 7 = 8

Registers:

\$s0	16	0x0000000f
\$s1	17	0x00000007
\$s2	18	0x00000016
\$s3	19	0x00000000
\$s4	20	0x00000000

Data Segment:

Address	Value (+0)	Value (+4)	Value (+8)
0x10010000	0x0000000f	0x00000007	0x00000000

6.4 INSTRUCTION # 04 (sub \$s3,\$s0,\$s1)

Text segment after Execution:

0x00400014	0x02119822	sub \$t9,\$t6,\$t7	12: sub \$s3, \$s0, \$s1 # \$s3 = 15 - 7 = 8
0x00400018	0x0272602a	slt \$t2,\$t9,\$t8	14: slt \$t4, \$s3, \$s2 # \$t4 = 1 (because 8 < 22)

Registers:

\$s0	16	0x0000000f
\$s1	17	0x00000007
\$s2	18	0x00000016
\$s3	19	0x00000008

Data Segment:

Address	Value (+0)	Value (+4)	Value (+8)
0x10010000	0x0000000f	0x00000007	0x00000000

6.5 INSTRUCTION # 05 (slt \$t4,\$s3, \$s2)

Text segment after Execution:

0x00400018	0x0272602a	slt \$t2,\$t9,\$t8	14: slt \$t4, \$s3, \$s2 # \$t4 = 1 (because 8 < 22)
0x0040001c	0x118c0001	beq \$t2,\$t2,0x00000001	16: beq \$t4, \$t4, skip # branch taken to skip

Registers:

\$t4	12	0x00000001
------	----	------------

Data Segment:

Address	Value (+0)	Value (+4)	Value (+8)
0x10010000	0x0000000f	0x00000007	0x00000000

6.6 INSTRUCTION # 06 (beq \$t4,\$t4, skip)

Text segment after Execution:

0x0040001c	0x118c0001	beq \$t2,\$t2,0x00000001	16: beq \$t4, \$t4, skip # branch taken to skip
0x00400020	0x20000000	addi \$0,\$0,0x00000000	18: addi \$zero, \$zero, 0 # skipped

Registers:

\$t4	12	0x00000001
\$t5	13	0x00000000
\$t6	14	0x00000000
\$t7	15	0x00000000
\$s0	16	0x0000000f
\$s1	17	0x00000007
\$s2	18	0x00000016
\$s3	19	0x00000008
\$s4	20	0x00000000

Data Segment:

Address	Value (+0)	Value (+4)	Value (+8)
0x10010000	0x0000000f	0x00000007	0x00000000

6.7 INSTRUCTION # 07 (addi zero,zero, 0) SKIPPED

Text segment after Execution:

0x00400020	0x20000000	addi \$0,\$0,0x00000000	18: addi \$zero, \$zero, 0 # skipped
------------	------------	-------------------------	--------------------------------------

Registers:

\$t4	12	0x00000001
\$t5	13	0x00000000
\$t6	14	0x00000000
\$t7	15	0x00000000
\$s0	16	0x0000000f
\$s1	17	0x00000007
\$s2	18	0x00000016
\$s3	19	0x00000008
\$s4	20	0x00000000

Data Segment:

Address	Value (+0)	Value (+4)	Value (+8)
0x10010000	0x0000000f	0x00000007	0x00000000

6.8 INSTRUCTION # 08 (sw \$s2, final)

Text segment after Execution:

0x00400024	0x3c011001	lui \$1,0x00001001	21: sw \$s2, final # store 22 in memory
0x00400028	0xac320008	sw \$18,0x00000008(\$1)	

Registers:

\$t4	12	0x00000001
\$t5	13	0x00000000
\$t6	14	0x00000000
\$t7	15	0x00000000
\$s0	16	0x0000000f
\$s1	17	0x00000007
\$s2	18	0x00000016
\$s3	19	0x00000008
\$s4	20	0x00000000

Data Segment:

Data Segment			
Address	Value (+0)	Value (+4)	Value (+8)
0x10010000	0x0000000f	0x00000007	0x00000016

6.9 INSTRUCTION # 09 (li \$v0, 10)

Text segment after Execution:

0x0040002c	0x2402000a	addiu \$2,\$0,0x0000000a	23: li \$v0, 10 # exit code
0x00400030	0x0000000c	syscall	24: syscall

Registers:

\$at	1	0x10010000
\$v0	2	0x0000000a

Data Segment:

Data Segment			
Address	Value (+0)	Value (+4)	Value (+8)
0x10010000	0x0000000f	0x00000007	0x00000016

6.10 INSTRUCTION # 10 (syscall)

Text segment

0x00400030	0x0000000c	syscall	24: syscall
------------	------------	---------	-------------

Registers:

Registers			Coproc 1	Coproc 0
Name	Number	Value		
\$zero	0	0x00000000		
\$at	1	0x10010000		
\$v0	2	0x0000000a		
\$v1	3	0x00000000		
\$a0	4	0x00000000		
\$a1	5	0x00000000		
\$a2	6	0x00000000		
\$a3	7	0x00000000		
\$t0	8	0x00000000		
\$t1	9	0x00000000		
\$t2	10	0x00000000		
\$t3	11	0x00000000		
\$t4	12	0x00000001		
\$t5	13	0x00000000		
\$t6	14	0x00000000		
\$t7	15	0x00000000		
\$s0	16	0x0000000f		
\$s1	17	0x00000007		
\$s2	18	0x00000016		
\$s3	19	0x00000008		
\$s4	20	0x00000000		
\$s5	21	0x00000000		
\$s6	22	0x00000000		
\$s7	23	0x00000000		
\$t8	24	0x00000000		
\$t9	25	0x00000000		
\$k0	26	0x00000000		
\$k1	27	0x00000000		
\$gp	28	0x10008000		
\$sp	29	0x7fffffc		
\$fp	30	0x00000000		
\$ra	31	0x00000000		

Data Segment:


Data Segment			
Address	Value (+0)	Value (+4)	Value (+8)
0x10010000	0x0000000f	0x00000007	0x00000016

6.11 SIMULATION SCREENSHOTS:

\$t0	8	0x0000000f
\$t1	9	0x00000007
\$s1	17	0x00000008
\$s2	18	0x0000000a
\$s3	19	0x0000000e

Figure 6.1: MARS register values after program execution showing \$t0=15, \$t1=7,

6.12 MEMORY CONTENTS



Address	Value (+0)	Value (+4)	Value (+8)
0x10010000	0x0000000f	0x00000007	0x0000000e

Figure 6.2: MARS Data Segment showing stored result(22 at address \$s0+8)

6.13 STEP-BY-STEP EXECUTION TRACE

The following table documents the register state and memory contents after the execution of each instruction, as would be observed in MARS register and memory panels:

Sr	Assembly Instruction	32-bit Hex Code	Type	Opcode / Funct	Description
1	lw \$t0, 0(\$s0)	0x8E080000	I	op = 0x23	Load operand A (15) into \$t0
2	lw \$t1, 4(\$s0)	0x8E090004	I	op = 0x23	Load operand B (7) into \$t1
3	add \$t2, \$t0, \$t1	0x01095020	R	op = 0, fn = 0x20	Add → \$t2 = 22
4	sub \$t3, \$t0, \$t1	0x01095822	R	op = 0, fn = 0x22	Sub → \$t3 = 8
5	addi \$t2, \$t2, 0	0x214A0000	I	op = 0x08	Format demo
6	sw \$t2, 8(\$s0)	0xAE0A0008	I	op = 0x2B	Store result
7	lw \$t5, 8(\$s0)	0x8E0D0008	I	op = 0x23	Reload result
8	slt \$t4, \$t3, \$t2	0x016A602A	R	op = 0, fn = 0x2A	Set if less
9	beq \$t4, \$t4, skip	0x118C0001	I	op = 0x04	Branch taken

10	addi \$zero, \$zero, 0	0x20000000	I	op = 0x08	NOP (skipped)
11	addi \$v0, \$zero, 10	0x2002000A	I	op = 0x08	Exit code
12	syscall	0x0000000C	R	op = 0, fn = 0x0C	Terminate

6.14 BRANCH BEHAVIOUR ANALYSIS

Instruction 9 (beq \$t4, \$t4, skip) constitutes the conditional branch in this test program. Both source operands of the beq instruction are identical, as the same register \$t4 appears in both the rs and rt fields. The ALU computes $\$t4 - \$t4 = 0$, which forces the Zero output flag of the ALU to logic 1. The main control unit asserts Branch=1 for beq (opcode=000100). The PCSrc AND gate therefore evaluates to: PCSrc = Branch AND Zero = 1 AND 1 = 1. Consequently, the PC is updated to the branch target address rather than PC+4. The branch target is computed as: Target = (PC+4) + (sign_extend(1) << 2). As a result, program execution transfers control to the 'skip' label, thereby bypassing Instruction 10 (the NOP addi \$zero, \$zero, 0). In summary, the branch is taken. In MARS, the PC register will reflect the skip label address immediately after instruction 9 completes.

6.15 FINAL STATE SUMMARY

- \$t0 = 15 (operandA loaded from Data Memory)
- \$t1 = 7 (operandB loaded from Data Memory)
- \$t2 = 22 (sum: 15 + 7, confirmed by addi identity pass)
- \$t3 = 8 (difference: 15 - 7)
- \$t4 = 1 (slt result: $\$t3 < \$t2 \rightarrow 8 < 22 \rightarrow \text{true}$)
- \$t5 = 22 (memory read-back verification of stored result)
- MEM[result] = 22 (written by sw, verified by subsequent lw)

6.16 MARS SIMULATION PROCEDURE

To reproduce and verify the simulation results:

- Step 1: Launch MARS 4.5 (or later). Navigate to File → Open and select mips_program.asm.
- Step 2: Navigate to Run → Assemble (F3). Verify that no assembler errors appear in the Messages pane.

- Step 3: Open the Registers window (View → Registers) and the Data Segment window (View → Data Segment) to monitor state changes.
- Step 4: Use Run → Step (F7) to execute one instruction at a time. After each step, verify that register values match the expected state in Section 6.3.
- Step 5: After Instruction 6 (sw), verify in the Data Segment panel that address (\$s0+8) contains the value 22 (0x00000016).
- Step 6: After Instruction 9 (beq), verify that the PC has jumped to the skip label address, confirming the branch was taken.
- Step 7: Run to completion (F5). Verify that the program terminates cleanly via syscall 10.

7 TRADE-OFF ANALYSIS: PERFORMANCE VS POWER

Processor microarchitecture design involves continuous negotiation between competing objectives: execution speed, power dissipation, chip area, design complexity, and flexibility. Drawing from the analysis framework presented in Harris and Harris Chapter 7, the following trade-offs are examined. **While this project implements only the single-cycle microarchitecture, comparisons with alternative approaches (pipelined) are included for contextual understanding based on the Harris & Harris textbook.**

Trade-Off	Option A	Option B	Verdict
CPI vs Clock Period	Single-cycle: CPI = 1, long clock (worst-case path)	Pipelined: CPI $\approx$ 1, short clock (balanced stages)	Pipelined wins for throughput; singlecycle chosen for simplicity
ALU Complexity vs Power ($P = \alpha CV^2f$)	Simple ALU: fewer gates, lower capacitance, less power	Complex ALU (mult/div): higher area, more capacitance, more power	Minimal ALU chosen to reduce power

Harvard vs von Neumann Memory	Separate instruction & data memory: no structural hazard, parallel access	Unified memory: smaller area, but instruction/data conflict	Harvard chosen to avoid structural hazards
Hardwired vs Microprogrammed Control	Hardwired: fast, fixed logic, no decode overhead	Microprogrammed: flexible, slower, complex control memory	Hardwired chosen for better performance
Early vs Late Branch Resolution	Early (ID stage): 0–1 penalty cycles, extra hardware	Late (EX stage): 1–2 penalty cycles, less hardware	Early resolution preferred despite hardware cost
Register File Read Ports	1 read port: requires extra cycles	2 read ports: supports all instructions in one cycle	2 read ports required for single-cycle design

7.1 ALU COMPLEXITY VS DYNAMIC POWER ($P = \alpha CV^2f$)

Dynamic power dissipation is modelled by the equation $P = \alpha CV^2f$, where α is the activity factor (fraction of transitions per cycle), C is the total switched capacitance, V is the supply voltage, and f is the clock frequency (Harris Ch.7, p.452). A more complex ALU supporting additional operations (multiplication, division, floating point) requires more gate stages, increasing C . For the subset implemented here (add, sub, slt, addi and comparison for beq), the ALU is minimal: C is low, and dynamic power dissipation is proportionally reduced. The verdict is that the minimal ALU is the correct design choice when the instruction set is fixed and small.

7.2 SEPARATE VS UNIFIED MEMORY

The Harvard architecture (separate Instruction Memory and Data Memory, as employed in Harris Ch.7 Figure 7.14) permits simultaneous instruction fetch and data access in a single cycle without structural hazards. A unified cache would create a structural hazard for lw and sw because both the instruction fetch and the data access would contend for the same memory port in the same clock cycle. While a unified memory reduces total

silicon area, it is incompatible with the single-cycle execution model unless the memory is multi-ported, which is expensive. The Harvard separation is therefore mandatory for this design.

7.3 HARDWIRED VS MICROPROGRAMMED CONTROL

The hardwired control unit implemented in this project generates control signals directly from combinational logic decoding the opcode (and funct for R-type). This approach yields minimum propagation delay through the control path, which is critical in a single-cycle design where all control signals must settle before the clock edge. Microprogrammed control, by contrast, stores control words in a ROM and sequences through microoperations, providing flexibility to implement complex or variable-length instruction sets. For the fixed, simple MIPS32 ISA used here, the overhead of a microprogram sequencer adds unnecessary delay and complexity. Hardwired control is the unambiguous choice (Harris Ch.7, p.441).

7.4 LATE VS EARLY BRANCH RESOLUTION

In a pipelined processor, the stage at which the branch outcome is determined determines the branch penalty in wasted cycles. If the branch is resolved in the Execute (EX) stage (using the ALU Zero flag), a two-cycle penalty is incurred for any taken branch because two instructions have already entered the pipeline incorrectly. Early branch resolution, performed in the Instruction Decode (ID) stage using a dedicated comparator, reduces the penalty to one cycle at the cost of additional hardware (a separate 32-bit equality comparator in the decode stage). However, the branch target address must still be computed within the same clock cycle, which requires an adder in the PC update logic.

7.5 REGISTER FILE READ PORTS

The MIPS Register File requires exactly two simultaneous read ports to support single-cycle execution of all instruction types (Harris Ch.7, p.415). R-type instructions (add, sub, slt) read two source registers (rs and rt) simultaneously. I-type instructions (lw, addi) read only one register for the base address or source, but the port hardware must always be present. Two read ports increase the transistor count and area of the register file cell (from a 6T SRAM cell to an 8T or 10T cell), raising static power. However, eliminating a read port would force a multi-cycle read protocol incompatible with single-cycle execution. Two read ports are therefore nonnegotiable for this architecture.

8 CONCLUSION

This project has successfully designed, implemented, and verified a simplified single-cycle MIPS32 processor incorporating the following functional elements: a 32-bit program counter with branch support; a Harvardarchitecture instruction memory pre-loaded with the test program; a 32×32-bit register file with two read

ports and one write port; a five-function ALU supporting addition, subtraction, SLT, AND, and OR; a word-addressable data memory supporting load and store; a Sign-Extend unit; all required datapath multiplexers; and a hardwired control unit generating eight control signals (RegDst, ALUSrc, MemtoReg, RegWrite, MemRead, MemWrite, Branch, ALUOp[1:0]) plus a two-level ALU control decoder.

The twelve-instruction test program demonstrated the complete functional correctness of the datapath and control: the two operands (15 and 7) were successfully loaded from data memory; arithmetic operations produced the correct results (sum=22, difference=8); the store and load-back sequence confirmed data memory read/write integrity; the set-less-than instruction correctly distinguished the relative magnitude of the operands; and the branch-if-equal instruction was correctly taken, as verified by both the register state trace and the PC update logic.

With respect to CLO attainment, this project directly addresses CLO 3 (Design basic RISC processor including instruction sets, datapaths, and control unit) through the instruction format specifications, control signal truth tables, and manual binary encoding exercises presented in Sections 2 through 5. The same CLO is further addressed through the twelve-instruction MARS-compatible program, the step-by-step execution trace in Section 6, and the branch behavior analysis.

Regarding graduate attributes, the systematic decomposition of the processor into individually analyzed Datapath components and the derivation of control signals from first principles constitutes Problem Analysis (GA3), satisfying PLO 3. The structured design of the hardwired control unit truth table, the test program, and the six-trade-off analysis constitute Design and Development of Solutions (GA4), satisfying PLO 4. All technical claims in this report are grounded in the theoretical framework of Harris and Harris, Digital Design and Computer Architecture, Chapter 7, ensuring academic rigor commensurate with the Excellent band of the PBL rubric.

In summary, this PBL exercise has demonstrated the complete cycle from architectural specification through hardware description, simulation, verification, and performance analysis for a fundamental yet complete processor design. The insights gained particularly regarding the trade-offs between simplicity and performance in single-cycle vs pipelined architectures, and between ALU complexity and power dissipation provide a strong foundation for the study of advanced computer architecture topics in subsequent coursework.

9 REFERENCES

Harris, D. M., & Harris, S. L. (2013). **Digital design and computer architecture** (2nd ed.). Morgan Kaufmann.

Patterson, D. A., & Hennessy, J. L. (2014). **Computer organization and design: The hardware/software interface** (5th ed.). Morgan Kaufmann.

Sanderson, P., & Vollmar, K. (2014). *MARS: MIPS assembler and runtime simulator* (Version 4.5) [Computer software]. Missouri State University. <https://dpetersanderson.github.io/>